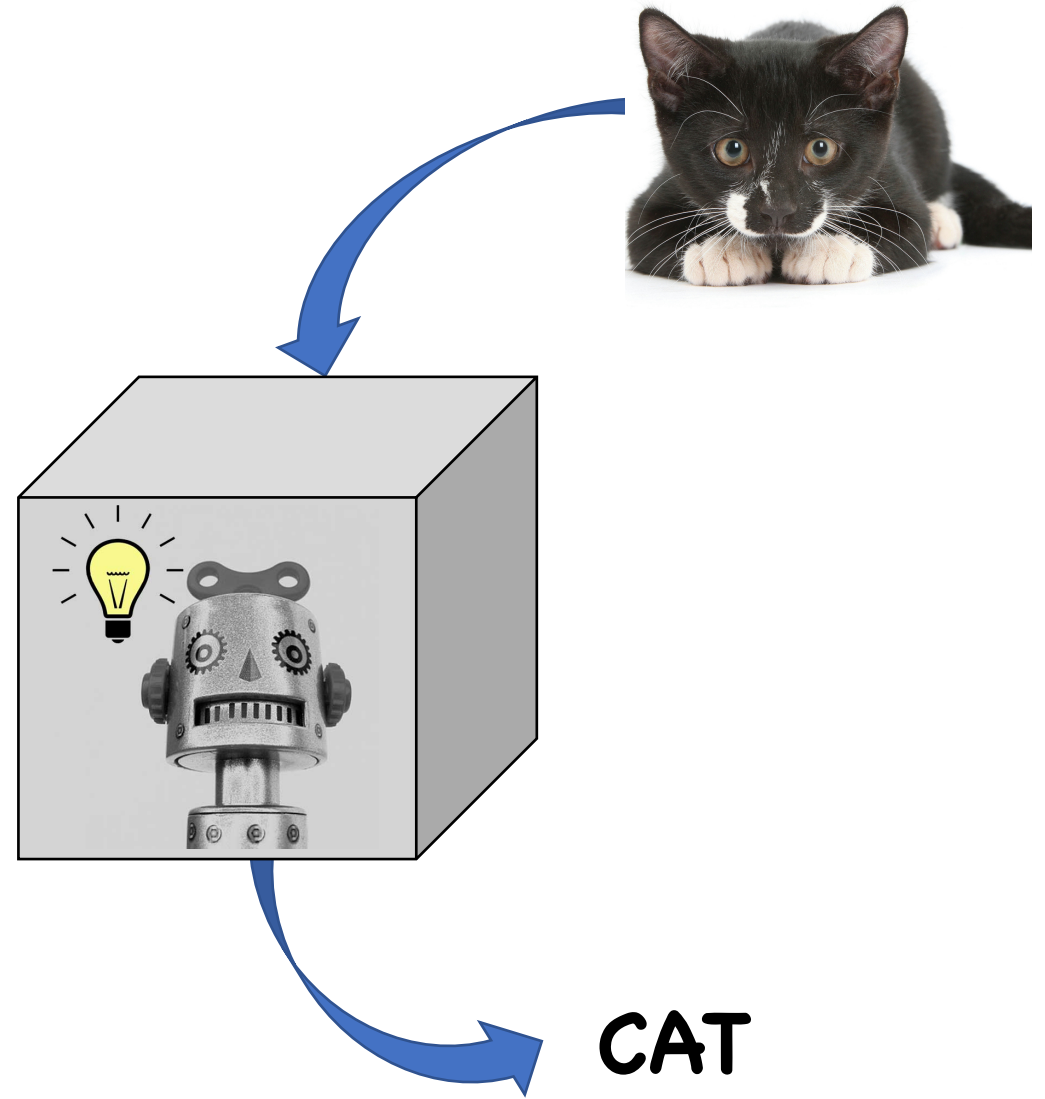



# Image Classification

Susmita Datta





# Dataset

## CIFAR-10

10 classes

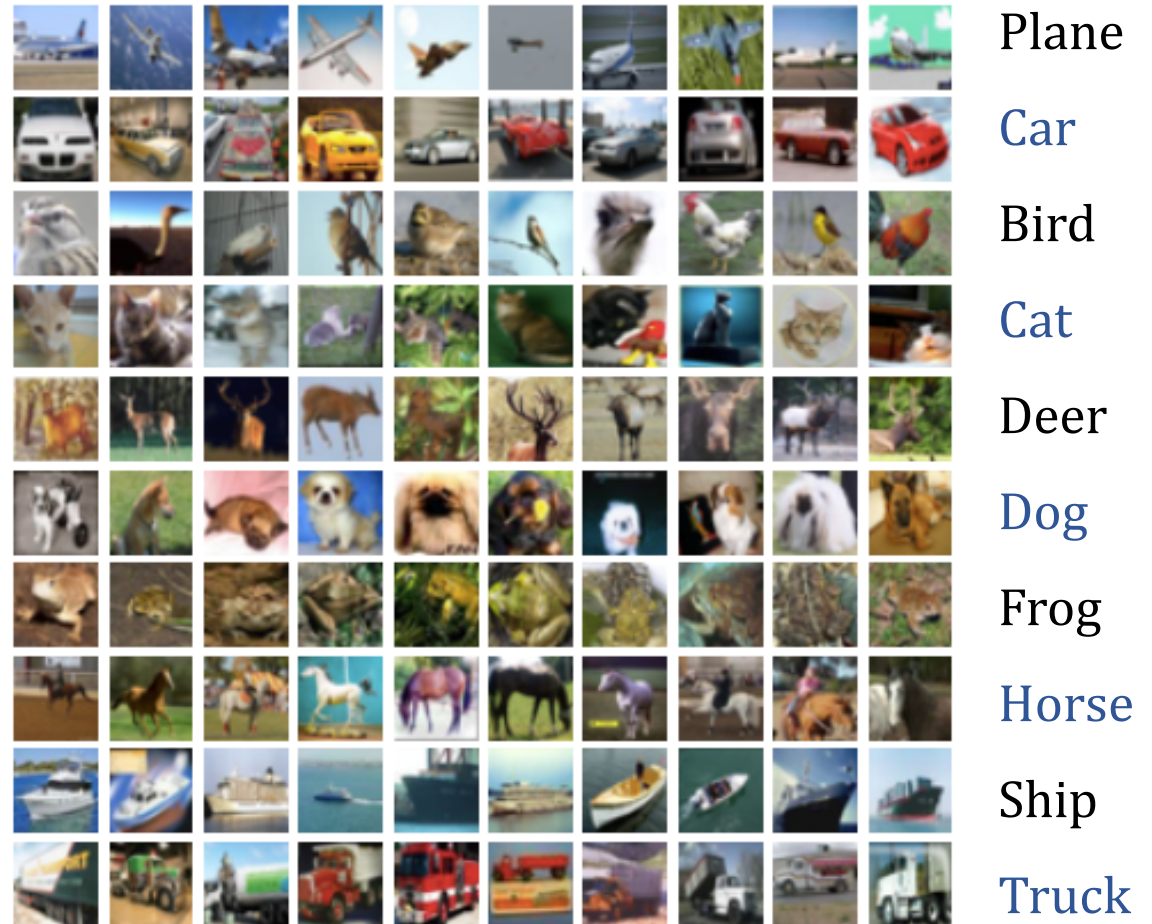
60,000 RGBs


- 50,000 training
- 10,000 test

32 pix × 32 pix × 3 pix

### Data source


<https://www.cs.toronto.edu/~kriz/cifar.html>





# Classification Models

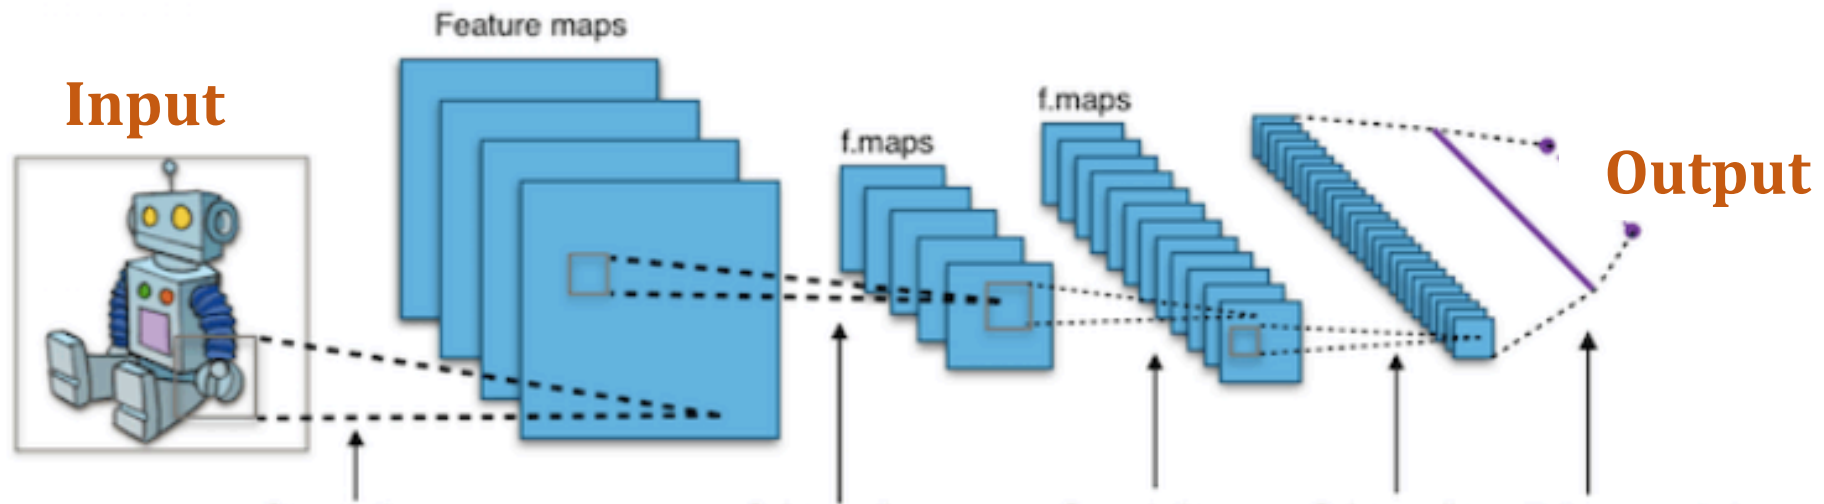
Accuracy	
Convolutional Neural Network	83%
K-nearest Neighbors	27%
Support Vector Machine	47%




# Convolutional Neural Network

## Convolutional Neural Network (ConvNet or CNN)

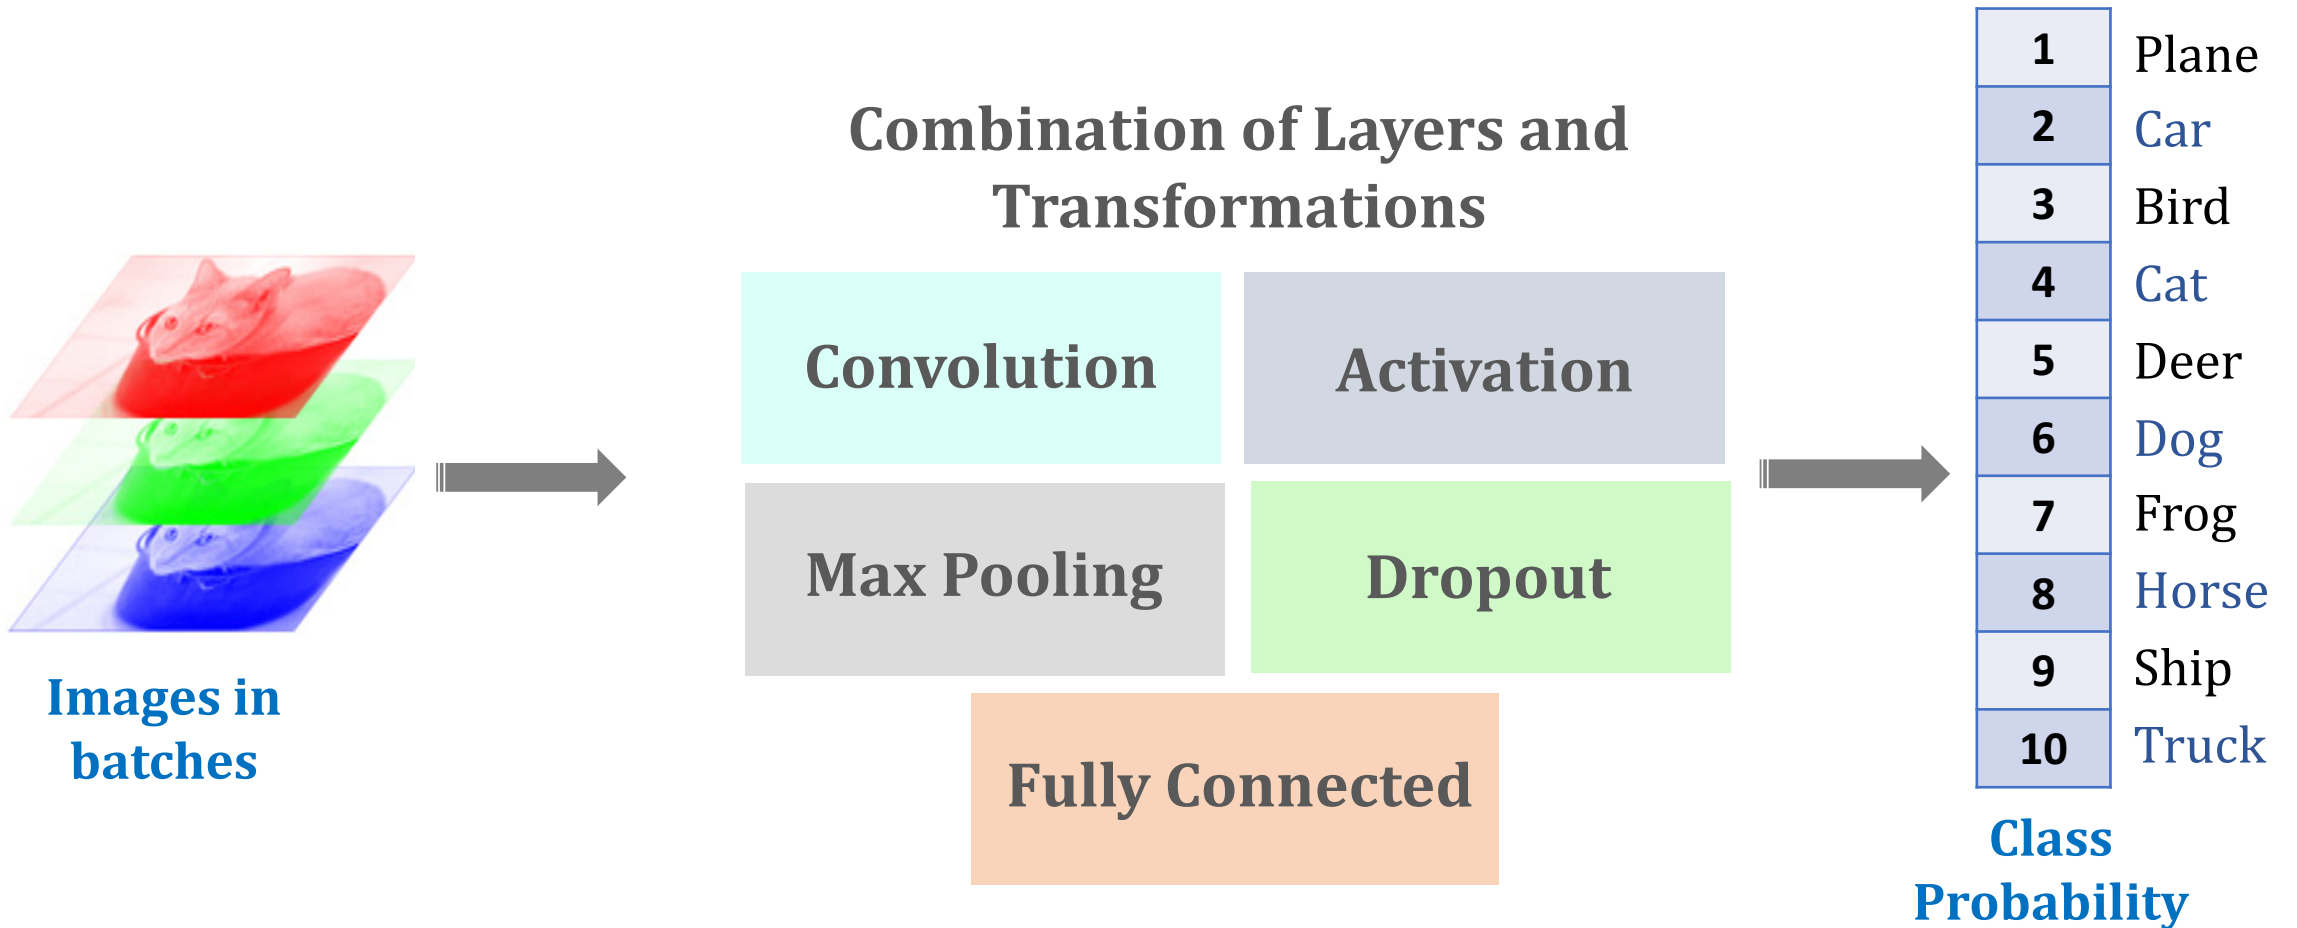
A deep learning process explicitly for images




Input goes through a series of transformations




# ConvNet Sequence Modeling






# Convolution Layer

Input	Transformation	Output
RGB image	Convolve images with learnable filters	Filter response (hidden layer)





# Convolution Layer

Input	Transformation	Output
RGB Images	Convolve images with learnable kernels	Kernel response





Image Pixels

Convolution applies weights to pixels



# Convolution Layer

Input	Transformation	Output
RGB Images	Convolve images with learnable kernels	Kernel response

A 3 x 3  
convolution  
kernel





Image Pixels

Convolution applies weights to pixels

Convolution kernel is a weight matrix





# Convolution Layer

Input	Transformation	Output
RGB Images	Convolve images with learnable kernels	Kernel response

Kernel is a  
weight matrix





Image Pixels



# Convolution Layer

Input	Transformation	Output
RGB Images	Convolve images with learnable kernels	Kernel response

Kernel is a  
weight matrix






Image Pixels

Weighted  
average



Convolution Response



# Convolution Layer

Input	Transformation	Output
RGB Images	Convolve images with learnable kernels	Kernel response

Kernel is a weight matrix






Image Pixels

Weighted average




Convolution Response




# Convolution Layer

Input	Transformation	Output
RGB Images	Convolve images with learnable kernels	Kernel response


Kernel is a  
weight matrix



Weighted  
average



Convolution Response



# Convolution Layer

Input	Transformation	Output
RGB Images	Convolve images with learnable kernels	Kernel response

Kernel is a  
weight matrix






Image Pixels

Weighted  
average



Convolution Response



# Convolution Layer

Input	Transformation	Output
RGB Images	Convolve images with learnable kernels	Kernel response

Kernel is a  
weight matrix






Image Pixels

Weighted  
average



Convolution Response



# Convolution Layer

Input	Transformation	Output
RGB Images	Convolve images with learnable kernels	Kernel response



**Original**


\*

1	0	-1
2	0	-2
1	0	-1

**Convolution  
kernel**






**Convolved**




# Activation Layer

Input	Transformation	Output
Output of Convolutional Layer	Apply ReLU function elementwise	Activation Response



Apply a threshold







# Max Pooling Layer

Input	Transformation	Output
Output of Previous Layer	Apply Max Pooling	Max-pooling image

Keep the maximum  
value of a region







# Max Pooling Layer

Input	Transformation	Output
Output of Previous Layer	Apply Max Pooling	Max pooling response

Keep the maximum  
value of a region







# Max Pooling Layer

Input	Transformation	Output
Output of Previous Layer	Apply Max Pooling	Max pooling response

Keep the maximum  
value of a region







# Max Pooling Layer

Input	Transformation	Output
Output of Previous Layer	Apply Max Pooling	Max pooling response


Keep the maximum  
value of a region






# Dropout Layer

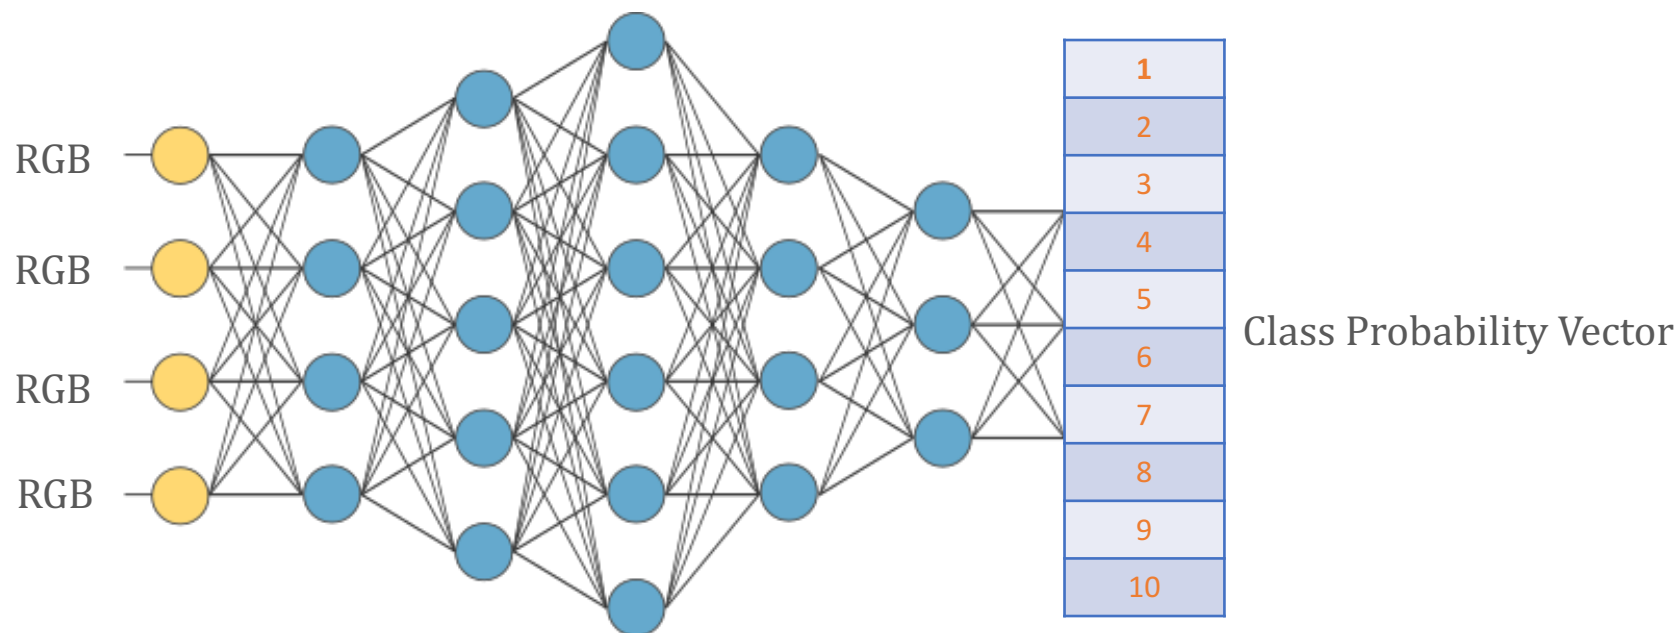
Randomly drop out neurons from a hidden layer







# Fully Connected Layer

Input	Transformation	Output
Output from the preceding layer	Mapping Matrix multiplication	N-dim probability vector (N = number of classes)







# ConvNet Model Summary





**Sequence of  
transformations  
in the model**




# ConvNet Model Summary



Image







# Training Set

## Image Augmentation



**Original**




**Rotation**



**Horizontal  
Flips**



**Vertical  
Flips**



# Increasing the Training Set

## Image Augmentation



**Original**



**Rotation**




**Horizontal  
Flips**




**Vertical  
Flips**

**Larger training set = original training set + augmented images**





# Accuracy and Loss

Model Accuracy




Model Loss






# Classification Accuracy






Overall accuracy: 83%




# Confusion Matrix


Actual Class






# Confusion Matrix





# Confusion Matrix



**Demo**